


Processing Near/In Memory

ISCA Tutorial (2019)

Website: <http://eyeriss.mit.edu/tutorial.html>


Processing Near vs. In Memory

Processing Near Memory



- data access energy and latency dominating
- data reuse and data compression

near memory




- computation still digital
- eliminates data transfer costs
- memory read energy dominates

Higher BW

Processing In Memory

deep in-memory



- memory access and computation combined
- mixed signal computation
- significant energy & latency reduction


Highest BW

Processing Near Memory

- Bring compute closer to the memory
- Benefits
 - Increase memory bandwidth
 - Reduce energy per access
- Challenges
 - Cost
- Embedded DRAM (eDRAM)
- 3D Stacked Memory (DRAM, SRAM)


eDRAM (DaDianNao)

- Advantages of eDRAM
 - 2.85x higher density than SRAM
 - 321x more energy-efficient than DRAM (DDR3)
- Store weights in eDRAM (36MB)
 - Target fully connected layers since dominated by weights




Stacked DRAM (NeuroCube)

- NeuroCube on Hybrid Memory Cube Logic Die
 - 6.25x higher BW than DDR3
 - HMC (16 ch x 10GB/s) > DDR3 BW (2 ch x 12.8GB/s)
 - Computation closer to memory (reduce energy)




Stacked DRAM (TETRIS)

- Explores the use of HMC with the Eyeriss **spatial architecture** and **row stationary dataflow**
- Allocates **more area to the computation (PE array)** than **on-chip memory (global buffer)** to exploit the low energy and high throughput properties of the **HMC**
 - **1.5x energy reduction, 4.1x higher throughput** vs. 2-D DRAM



Stacked SRAM (Quest)

- Explores use of stacked SRAM which has lower latency and power dissipation than DRAM
 - Lower storage density than DRAM but can stack more due to low power density
- Use inductive coupling to connect to SRAM rather than TSV for lower integration cost
 - $9.6\text{Gb/s per channel} \times 24 \text{ channels} = 28.8 \text{ GB/s}$




Processing In Memory

- Embed processing in peripheral or perform compute in analog domain
- Benefits
 - read multiple weights in parallel (high bandwidth)
 - perform multiple computations in parallel (high throughput)
 - Increase density of compute engines (lower capacitance on input delivery)
 - reduce number of conversions (sense amp only applied to final accumulated value)
- Challenges
 - non-idealities of analog compute (non-linearity, sensitivity to process and temperature variations)
 - conversion cost from analog to digital (fewer conversions, but each one is more expensive)
 - increase size of bit cell and increase size in peripheral (reduce storage density)
 - limited precision that can be stored in each storage element (e.g. bit cell, device); need to use bit serial or multiple elements if want to support higher precision; also may need to use two arrays to enable differential signal to represent negative values
- Memory Technology
 - SRAM, DRAM, NVM (Flash, memristors)

Dataflow for PIM




- Weight stationary dataflow



- Overhead in conversion from analog to digital and digital back to analog

DRAM (DRISA)

- Build multiplier and adders from NOR logic in DRAM
 - Add shift logic to output of DRAM before write back
- Requires multiple cycles
 - Number of cycles depends on bit width
 - Tradeoff parallelism for cycles
- Parallelism dictated by array width and number of arrays
- Demonstrate on BNN: 1b weights, 8 bit activation for AlexNet, VGG-16, VGG-19, ResNet




Analog Computation

- Conductance = Weight
- Voltage = Input
- Current = Voltage × Conductance
- Sum currents for addition

$$Output = \sum Weight \times Input$$

Input = V_1, V_2, \dots


Filter Weights = G_1, G_2, \dots (conductance)




Weight Stationary Dataflow

Compute Dot Product in SRAM


Approaches vary in terms of how data applied to *BL* and/or *WL*



[J. Zhang,
VLSI 2016]




[S. Gonugondla,
ISSCC 2018]



Bit-cell discharge current:
 $I_{cell} \propto (V_{WL} - V_t)$
 $V_a \propto I_{cell}$

Conventional [4]

V_a varies widely
due to I_{cell} variation



*Assumption:
Both DACs are ideal

Directly apply V_a on
GBL(LBL) by DAC

$$V_a \propto I_{cell}$$

Proposed

[A. Biswas, ISSCC 2018]

V_a has no variation due to I_{cell}

Analog Compute with Flash Memory

Floating Gate: Program by changing threshold voltage (shift I-V curve)





Image source:
Diorio, ISFIIIE 1999

	Enterprise With DRAM	Enterprise No-DRAM	Edge With DRAM	Edge No-DRAM	Mythic NVM
SRAM	<50 MB	100+ MB	< 5 MB	< 5 MB	< 5 MB
DRAM	8+ GB	-	4-8 GB	-	-
Power	70+ W	70+ W	3-5 W	1-3 W	1-5 W
Sparsity	Light	Light	Moderate	Heavy	None
Precision	32f / 16f / 8i	32f / 16f / 8i	8i	1-8i	1-8i
Accuracy	Great	Great	Moderate	Poor	Great
Performance	High	High	Very Low	Very Low	High
Efficiency	25 pJ/MAC	2 pJ/MAC	10 pJ/MAC	5 pJ/MAC	0.5 pJ/MAC


Source: Mythic,
Hot Chips 2018

Memristor Computation


Use memristors as programmable weights (resistance)

- **Advantages**
 - **High Density (< 10nm x 10nm size*)**
 - ~30x smaller than SRAM**
 - 1.5x smaller than DRAM**
 - **Non-Volatile**
 - **Operates at low voltage**
 - **Computation within memory (in situ)**
 - Reduce data movement


Memristor



(a) Conceptual view
of a ReRAM cell




(b) I-V curve of bipolar
switching




(c) schematic view of a
crossbar architecture

$$b_j = \sigma\left(\sum_{\forall i} a_i \cdot w_{i,j}\right)$$



(a) An ANN with one input
and one output layer





(b) using a ReRAM crossbar
array for neural computation

Challenges with Memristors


- Limited Precision
- A/D and D/A Conversion
- Array Size and Routing
 - Wire dominates energy for array size of $1k \times 1k$
 - IR drop along wire can degrade read accuracy
- Write/programming energy
 - Multiple pulses can be costly
- Variations & Yield
 - Device-to-device, cycle-to-cycle
 - Non-linear conductance across range

ISAAC

- eDRAM using memristors
- 16-bit dot-product operation
 - 8 x 2-bits per memristors
 - 1-bit per cycle computation
 - Trade off area and cycles to address low precision



ISAAC




Eight 128x128 arrays per IMA

12 IMAs per Tile

14x12 Tiles in ISAAC



PRIME

- Bit precision for each 256x256 ReRAM array
 - 3-bit input, 4-bit weight (2x for 6-bit input and 8-bit weight)
 - Dynamic fixed point (6-bit output)
- Reconfigurable to be main memory or accelerator
 - 4-bit MLC computation; 1-bit SLC for storage



Fabricated Arrays

- Ternary weights and binary activations
- Use two arrays for ternary weights {-1, 0, +1}
- [Su, VLSI 2017]
 - Array size: 32x32 (2 arrays)
 - Accumulate 32 values into a 3-bit output
 - Demo Task
 - Input: 5x5 image (25 values)
 - Classify with 2 FC (25x32x4)
 - Output: 4 directions
- nvCIM [Chen, ISSCC 2018]
 - Array size: 512 x 256 (8 arrays)
 - Accumulate 512 values into a 3-bit output
 - Max 32 rows activated at a time due to sensing margin and yield
 - Demo Task
 - Input: 5x5 image (MNIST) – 96.2% accuracy
 - Classify with 2 FC (25x25x3) ??
 - Output: 9 digits



Optical Neural Network

Matrix Multiplication in the Optical Domain

The photodetection rate is 100 GHz

"In principle, such a system can be at least two orders of magnitude faster than electronic neural networks (which are restricted to a GHz clock rate)"

